

10BASE-T1S through a media converter, measured end to end by two ESP32-S3 boards

Run 2026-10-06 15:27:03 — generated by `make_duo_report.py` from `duo_suite_20261006_153816.json`.

Abstract

A LAN8651 MAC-PHY node on an ESP32-S3 and a second ESP32-S3 on 100BASE-TX exchanged UDP through a 100BASE-TX/10BASE-T1S converter, with no PC in the data path. Both boards ran the same firmware and measured each other. One direction at a time, the bus carried **8.51 Mbit/s onto T1S and 7.97 Mbit/s off it with no datagram lost** (sequence-checked), 87 % and 81 % of the 9.80 Mbit/s the 10BASE-T1S frame model allows. A 64-byte UDP round trip took **3.08 ms median** (p99 3.66 ms); 1472 bytes took 8.40 ms. Over 60 s at 8 Mbit/s each way, 7 of 39631 (tx->node) and 0 of 39400 (node->tx) datagrams were lost. With both directions loaded at once, the node's stream arrived whole while the converter's own stream lost most of its datagrams — the converter, not the bus, is the bottleneck there. Zenoh ran on the same two boards **peer to peer over UDP multicast, with no router**: a pub/sub round trip took 5.2–5.3 ms median in either direction, and bulk puts delivered up to 4.9 Mbit/s of payload.

1. Setup



ESP-A: ESP32-S3 + LAN8651 HAT converter ESP-B: ESP32-S3 + W5500 SPI 20.00 MHz, PLCA 0 of 2 (coord.) LAN8670 + LAN9355 W5500 SPI 40 MHz, polled 1 ms 192.168.100.65 = 10BASE-T1S = [T1S TX] — 100BASE-TX — 192.168.100.66	
node	LAN8651 rev 2, firmware <code>t1s_node</code> mode node, UDP echo on port 7, sink on port 9
peer	W5500 (polled every 1 ms by the IDF driver), <code>t1s_node</code> mode tx, same echo and sink
converter	3-port 10BASE-T1S / 100BASE-T1 / 100BASE-TX media converter (LAN8670 PHY, LAN9355 switch); dials ID 1, count 0, so the HAT is the PLCA coordinator
traffic	<code>blast</code> : UDP with a 32-bit sequence number in the first 4 bytes, optionally paced to a target rate; <code>sink</code> : counts, checks sequence (loss, reordering, duplicates) and histograms inter-arrival gaps (20 µs bins)
latency	<code>rtt</code> : UDP echo, one datagram in flight, timed with <code>esp_timer</code> (1 µs), 3 ms between probes, 100 ms timeout
rate accounting	Mbit/s counts UDP payload + 42 B of Ethernet/IP/UDP headers per datagram, as both tools report

T1S model. Each datagram of L payload bytes occupies (L + 42 + 4 FCS + 8 preamble + 12 IFG) bytes on the wire, plus about 52 bit times of PLCA overhead per frame (the idle node's 32-bit transmit opportunity and a share of the BEACON, burst mode off). At 1472 B this gives a ceiling of **9.80 Mbit/s** in the accounting above.

2. Results

2.1 Round-trip time


Fig. 1 Round-trip time across converter + 10BASE-T1S, by payload size
measured on ESP-B (W5500)



direction	payload	n	lost	min	median	p99	max	σ
ESP-B → HAT (onto T1S)	64 B	400	0	2.45	3.08	3.66	4.25	0.285
ESP-B → HAT (onto T1S)	256 B	400	0	3.14	3.71	4.46	4.91	0.347
ESP-B → HAT (onto T1S)	512 B	400	0	4.15	4.66	5.33	5.45	0.422
ESP-B → HAT (onto T1S)	1024 B	400	0	6.26	6.36	7.41	7.86	0.318
ESP-B → HAT (onto T1S)	1472 B	400	0	7.81	8.40	9.14	9.34	0.389
HAT → ESP-B (off T1S)	64 B	400	0	2.59	2.71	3.62	3.85	0.160
HAT → ESP-B (off T1S)	256 B	400	0	3.41	4.16	4.57	4.66	0.112
HAT → ESP-B (off T1S)	512 B	400	0	4.62	4.74	5.14	5.62	0.115
HAT → ESP-B (off T1S)	1024 B	400	0	6.72	6.87	7.46	8.77	0.162
HAT → ESP-B (off T1S)	1472 B	400	0	8.24	8.77	9.42	9.77	0.148

All times in ms.

Fig. 2 RTT vs payload, linear fit (dotted)




- ESP-B → HAT (onto T1S): $RTT \approx 2.76 \text{ ms} + 3.73 \mu\text{s} \times \text{payload bytes}$.
- HAT → ESP-B (off T1S): $RTT \approx 2.72 \text{ ms} + 4.11 \mu\text{s} \times \text{payload bytes}$.

Millisecond steps. Seen from ESP-B, round trips cluster at whole milliseconds: 38 % of the 64 B samples lie within 0.1 ms of an integer number of ms, and larger payloads jump from one step to the next (Fig. 1, left). That is the W5500 driver's 1 ms receive poll: a reply that reaches the chip waits for the next poll tick. Seen from the HAT, whose LAN8651 interrupts on receive, the same paths spread continuously (Fig. 1, right).

Reading it. Every payload byte crosses each link twice per round trip. Serialisation alone predicts 10BASE-T1S (0.8 $\mu\text{s/B}$) → 1.60, 100BASE-TX (0.08 $\mu\text{s/B}$) → 0.16, LAN8651 SPI 20.00 MHz (0.40 $\mu\text{s/B}$) → 0.80, W5500 SPI 40 MHz (0.2 $\mu\text{s/B}$) → 0.40 μs , **2.96 $\mu\text{s/B}$ in total; the fit gives 3.92 $\mu\text{s/B}$.** The remaining 0.96 $\mu\text{s/B}$ is per-byte work in software (copies through lwIP and both drivers). The intercept (≈ 2.7 ms) is fixed per-packet cost: the W5500 driver's 1 ms receive poll on ESP-B, lwIP and echo-task turnarounds, and PLCA transmit opportunities. The PC-based run of 2026-10-01 had no polling on its side and measured 0.85 ms for 64 B.

2.2 Throughput, one direction at a time


Fig. 3 Load sweep, one direction at a time (paced; right-most point unpaced; hollow = sender saturated, x is the asked rate)



direction	offered	delivered	sent	received	loss	reordered	gap p50 / p99 (μs)
ESP-B → HAT (onto T1S)	1.00	1.00	332	332	0.00 %	0	12500 / 12500
ESP-B → HAT (onto T1S)	2.00	2.00	662	662	0.00 %	0	6500 / 6500
ESP-B → HAT (onto T1S)	3.00	3.00	992	992	0.00 %	0	4500 / 4500
ESP-B → HAT (onto T1S)	4.00	4.00	1322	1322	0.00 %	0	3500 / 3500
ESP-B → HAT (onto T1S)	5.00	5.00	1652	1652	0.00 %	0	2500 / 2500
ESP-B → HAT (onto T1S)	6.00	6.00	1983	1983	0.00 %	0	2500 / 2500
ESP-B → HAT (onto T1S)	7.00	7.00	2313	2313	0.00 %	0	1730 / 2500
ESP-B → HAT (onto T1S)	8.00	8.00	2642	2642	0.00 %	0	1490 / 2500
ESP-B → HAT (onto T1S)	8.51	8.51	2810	2810	0.00 %	0	1410 / 1950
ESP-B → HAT (onto T1S)	9.01	8.97	2974	2964	0.34 %	0	1330 / 1750
ESP-B → HAT (onto T1S)	9.51	8.96	3140	2961	5.70 %	0	1330 / 1730
ESP-B → HAT (onto T1S)	10.00	9.00	3303	1531	53.65 %	0	1330 / 1690
ESP-B → HAT (onto T1S)	13.48	9.00	4451	1245	72.03 %	0	1330 / 1690
HAT → ESP-B (off T1S)	1.00	1.00	332	332	0.00 %	0	12500 / 13336
HAT → ESP-B (off T1S)	2.00	2.00	662	662	0.00 %	0	6500 / 7367
HAT → ESP-B (off T1S)	3.00	3.00	992	992	0.00 %	0	3500 / 5500
HAT → ESP-B (off T1S)	4.00	4.00	1322	1322	0.00 %	0	2500 / 4175
HAT → ESP-B (off T1S)	5.00	5.00	1651	1651	0.00 %	0	2500 / 3500
HAT → ESP-B (off T1S)	6.00	6.00	1983	1983	0.00 %	0	1990 / 3380
HAT → ESP-B (off T1S)	6.99	7.00	2310	2310	0.00 %	0	1910 / 2500
HAT → ESP-B (off T1S)	7.94	7.94	2623	2623	0.00 %	0	1730 / 2500
HAT → ESP-B (off T1S)	8.5	7.95	2624	2624	0.00 %	0	1730 / 2500
HAT → ESP-B (off T1S)	9	7.96	2629	2629	0.00 %	0	1710 / 2500
HAT → ESP-B (off T1S)	9.5	7.96	2628	2628	0.00 %	0	1710 / 2500
HAT → ESP-B (off T1S)	10	7.97	2630	2630	0.00 %	0	1710 / 2500
HAT → ESP-B (off T1S)	7.96	7.96	2629	2629	0.00 %	0	1710 / 2500


Overload stalls the stream. In 2 row(s) the receiver's window is shorter than the blast: ESP-B → HAT (onto T1S) at 10.00 Mbit/s: 2.1 of 4.0 s, seq gaps 162; ESP-B → HAT (onto T1S) at 13.48 Mbit/s: 1.7 of 4.0 s, seq gaps 610. Delivery ran at the bus ceiling and then **stopped** for the rest of the blast, which is where the large loss figures at and above ~10 Mbit/s come from; the sequence gaps inside the window are the smaller part. The node received normally again in the next run. Overloading the converter's T1S side therefore costs more than the excess: the segment goes quiet for seconds. §2.2.1 locates the stop.

Fig. 7 Arrival spacing (2 ms histogram cap)



2.3 Payload size

Fig. 4 Unpaced blasts by payload size




direction	payload	delivered (Mbit/s)	frames/s	model ceiling (Mbit/s)	loss
ESP-B → HAT (onto T1S)	64 B	1.96	2306	7.77	0.0 %
ESP-B → HAT (onto T1S)	128 B	3.03	2226	8.48	0.0 %
ESP-B → HAT (onto T1S)	256 B	4.93	2068	9.07	0.0 %
ESP-B → HAT (onto T1S)	512 B	7.57	1708	9.48	5.8 %
ESP-B → HAT (onto T1S)	1024 B	8.58	1007	9.72	87.2 %
ESP-B → HAT (onto T1S)	1472 B	8.98	742	9.80	86.5 %
HAT → ESP-B (off T1S)	64 B	2.25	2655	7.77	0.0 %
HAT → ESP-B (off T1S)	128 B	3.18	2336	8.48	0.0 %
HAT → ESP-B (off T1S)	256 B	4.47	1875	9.07	0.0 %
HAT → ESP-B (off T1S)	512 B	5.96	1345	9.48	0.0 %
HAT → ESP-B (off T1S)	1024 B	7.35	862	9.72	0.0 %
HAT → ESP-B (off T1S)	1472 B	7.94	656	9.80	0.0 %

Small datagrams are bounded by per-packet cost in the boards (lwIP, SPI transactions, the W5500 poll), not by the bus: the frame rate stays far below what 10BASE-T1S could carry. Unpaced in the ESP-B → HAT direction, the W5500 offers more than the bus can take and the converter drops the excess.

2.4 Both directions at once

Fig. 5 Both directions at once




offered each	ESP-B → HAT delivered	loss	HAT → ESP-B delivered	loss
1	1.00	0.0 %	1.00	0.0 %
2	2.00	0.0 %	2.00	0.0 %
3	1.13	62.4 %	3.00	0.0 %
4	3.73	6.7 %	4.00	0.0 %
4.5	1.91	57.7 %	4.50	0.0 %
5	1.71	65.7 %	4.87	0.0 %

The HAT's stream arrives whole at every level. The stream the converter must put onto T1S loses most of its datagrams as soon as both directions are busy, even at 1–2 Mbit/s each, far below the bus capacity. That is the converter's T1S side failing to win transmit opportunities while it is receiving, matching its vendor's own issue 4.3 ("bidirectional traffic through the converter does not work properly — LAN8670"), and the 2026-10-01 PLCA run with a PC in ESP-B's place.

2.5 Latency under load

Fig. 6 Latency under reverse load



HAT → ESP-B load	probes	lost	median	p99
0 Mbit/s	300	0	3.08	3.59
3 Mbit/s	300	5	4.65	5.75
6 Mbit/s	300	67	3.11	6.71
9 Mbit/s	300	69	3.09	8.58

Echo probes from ESP-B must cross the converter onto the bus the HAT is streaming on, so they meet the same converter limit as §2.4: answered probes keep their latency, but a growing share is lost.


2.6 Soak

direction	duration	offered	delivered	sent	received	lost	reordered	duplicates
ESP-B → HAT (onto T1S)	60 s	8.0	8.00	39631	39624	7	0	0
HAT → ESP-B (off T1S)	60 s	7.95	7.95	39400	39400	0	0	0

2.7 Zenoh, peer to peer with no router

Both boards ran zenoh-pico in **peer mode over UDP multicast** (udp/224.0.0.224:7447): no zenohd, no PC. Each board echoes the other's test/ping/<node> on test/pong/<node>; the pinging board timestamps with esp_timer. RTT was taken at 50 Hz with the boards' normal Zenoh traffic running (20 Hz signal, 2 Hz hello, 1 Hz stats each); bulk transfers ran with that traffic paused.


Fig. 8 Zenoh ping/pong vs raw UDP echo



pinged from	pongs	min	median	p99	max	raw UDP 64 B median	Zenoh overhead
ESP-B	589 of ~600	4.88	5.27	6.63	6.99	3.08	+2.19 ms
HAT	591 of ~600	4.95	5.18	6.44	7.18	2.71	+2.46 ms

All times in ms. *pongs of ~N*: the board keeps its last 600 round trips, and pings sent while the previous pong was outstanding are not all answered (each board echoes from a short queue).

Fig. 9 Zenoh put/subscribe bulk transfer, peer over UDP multicast, no router



direction	payload	put (msg/s)	received (msg/s)	put (Mbit/s)	received (Mbit/s)	lost
ESP-B → HAT (onto T1S)	32 B	944	944	0.24	0.24	0.0 %
ESP-B → HAT (onto T1S)	128 B	919	920	0.94	0.94	0.0 %
ESP-B → HAT (onto T1S)	256 B	887	887	1.82	1.82	0.0 %
ESP-B → HAT (onto T1S)	512 B	833	782	3.41	3.20	6.1 %
ESP-B → HAT (onto T1S)	1024 B	728	596	5.96	4.88	18.1 %
HAT → ESP-B (off T1S)	32 B	1129	910	0.29	0.23	16.0 %
HAT → ESP-B (off T1S)	128 B	1017	889	1.04	0.91	10.6 %
HAT → ESP-B (off T1S)	256 B	912	860	1.87	1.76	4.5 %
HAT → ESP-B (off T1S)	512 B	775†	775	3.17†	3.17	0.0 %
HAT → ESP-B (off T1S)	1024 B	585	585	4.79	4.79	0.0 %

† the publisher's own report line did not reach the console; the put count is the subscriber's sequence range (exact unless the final messages were lost), over the blast duration of the other rows.

What limits it. Puts are unpaced, so each row is what zenoh-pico sustains. For small payloads the rate is set first by the **publisher** (put / received: ESP-B → HAT (onto T1S) 32 B: 944 / 944, ESP-B → HAT (onto T1S) 128 B: 919 / 920, HAT → ESP-B (off T1S) 32 B: 1129 / 910, HAT → ESP-B (off T1S) 128 B: 1017 / 889; the same boards receive 2400+ raw UDP frames/s, §2.3). Unbatched, every put is its own UDP datagram and a put call takes 808–1632 μ s on the sender, which holds small messages near 1000/s. Where the subscriber kept up (ESP-B → HAT (onto T1S) 32 B, ESP-B → HAT (onto T1S) 128 B) the publisher is the whole limit; where it did not (HAT → ESP-B (off T1S) 32 B, HAT → ESP-B (off T1S) 128 B) the receiving board also drops some. Large payloads from ESP-B lose datagrams at an *average* rate the bus carried cleanly as paced UDP (§2.2): unpaced puts leave the W5500 back to back at 100 Mbit/s, so the converter's buffer meets bursts well above 10 Mbit/s (our reading; the converter exposes no drop counter). In the other direction the HAT's SPI/TC6 path paces its puts.

With batching (zenoh blast ... batch: zp_batch_start/stop around the run, puts packed into datagrams until one is full):

direction	payload	unbatched put (msg/s)	batched put (msg/s)	put call	batched received (msg/s)	received (Mbit/s)	lost
ESP-B → HAT (onto T1S)	32 B	944	7483	99 μs	6416	1.64	14.0 %
ESP-B → HAT (onto T1S)	128 B	919	4395	192 μs	3864	3.96	12.1 %
ESP-B → HAT (onto T1S)	256 B	887	2745	321 μs	2295	4.70	16.4 %
ESP-B → HAT (onto T1S)	512 B	833	1331	687 μs	1148	4.70	13.7 %
ESP-B → HAT (onto T1S)	1024 B	728	716	1308 μs	595	4.87	16.9 %
HAT → ESP-B (off T1S)	32 B	1129	7430	105 μs	7282	1.86	1.7 %
HAT → ESP-B (off T1S)	128 B	1017	3661	240 μs	3663	3.75	0.0 %
HAT → ESP-B (off T1S)	256 B	912	2242	409 μs	2243	4.59	0.0 %
HAT → ESP-B (off T1S)	512 B	775	1107	851 μs	1108	4.54	0.0 %
HAT → ESP-B (off T1S)	1024 B	585	578	1652 μs	578	4.74	0.0 %

Batching lifts small messages several-fold because the per-datagram cost is paid once per batch; a payload that fills a datagram by itself gains nothing. The price is latency: a batched put waits until the batch is flushed, so this is for bulk, not for control messages.

3. Discussion

- **The T1S segment is not the limit for one-way traffic.** Paced at up to 9 Mbit/s, neither direction lost a datagram, and the unpaced HAT → ESP-B stream sits close to the model ceiling.
- **Do not offer more than the bus.** Above the ~9.8 Mbit/s ceiling the excess is not simply dropped: the stream stalls for seconds, and the stall is in the converter, not the node (§2.2.1). A T1S edge that can be overloaded from a faster segment needs shaping at the entry (the bridge firmware's job), not just a big buffer.
- **The converter is the limit for two-way traffic.** Loss in §2.4 appears at a few Mbit/s total, so it is not capacity; it is the converter's own T1S transmitter. A PLCA-aware MAC-PHY on both ends (or a converter that reserves its transmit opportunities) would remove it; that is why the HAT keeps a LAN8651 rather than a PHY behind a switch.
- **Zenoh needs no infrastructure on this segment.** Two microcontrollers formed a Zenoh network on their own over T1S + converter, with multicast discovery and no router; its cost over raw UDP is a few milliseconds of round trip and a ~1000 msg/s receive ceiling in zenoh-pico on this platform.
- **Board latency dominates RTT.** At 64 B the bus contributes well under 0.2 ms of a ~3 ms round trip; the W5500's 1 ms polling and per-packet software dominate. An interrupt-driven W5500 (the T-ETH-Elite does not route its INT line) or a second LAN8651 node would cut it.

4. Limitations

- One converter unit, one bus length (~1 m), two nodes; no EMC or cable-length effects are captured.
- Rates are application-level UDP; preamble, FCS and IFG are not counted (the model curve accounts for them).
- Arrival gaps above 2 ms fall into one histogram bin, so the p99 for slow rates is a lower bound.
- sink computes its rate over first-to-last arrival, which reads ~0.5 % high against the sender's window; that is why the 1472 B points can sit a hair above the model ceiling.

- RTT probes are sequential (one in flight); they measure latency, not latency under pipelining.
- rtt's 3 ms probe spacing is not synchronised with PLCA cycles; queueing behind the other node's transmit opportunity is part of the measured distribution, by design.

5. Reproduce

Firmware: elite-tls-hat/firmware/tls_node (both boards); console: esp32-tls-bridge/pc/tls_console/hat_server.py, POST /api/suite (or the Two ESPs tab); report: python3 pc/tls_console/make_duo_report.py <json>.